


Writing Models in Keras

ML Bootcamp India (2022)



Sayak Paul
ML Engineer at Carted
[@RisingSayak](#)




\$whoami

- ML Engineer at [Carted](#)
- Open-source 🥑 (Keras, KerasCV, 🤗 Transformers etc.)
- Netflix nerd
- Coordinates at [sayak.dev](#)

Agenda

- Several ways to write models in Keras
- When to use what
- Code
- QnA


Materials are here:
bit.ly/code-mlp-2022




Two broad styles: Symbolic and Imperative

- Symbolic:
 - Your model is sequential stack of layers (Sequential)
 - Your model is basically a DAG (Functional)

Two broad styles: Symbolic and Imperative



DAG



Stack

Credits: [What are Symbolic and Imperative APIs in TensorFlow 2.0?](#)

Two broad styles: Symbolic and Imperative

- Symbolic
- Imperative:
 - You need to have granular controls over how each component in your model is executed.

Two broad styles: Symbolic and Imperative

```
class CNN_Encoder(tf.keras.Model):  
    def __init__(self, embedding_dim):  
        super(CNN_Encoder, self).__init__()  
        self.fc = tf.keras.layers.Dense(embedding_dim)  
  
    def call(self, x):  
        x = self.fc(x)  
        x = tf.nn.relu(x)  
        return x
```

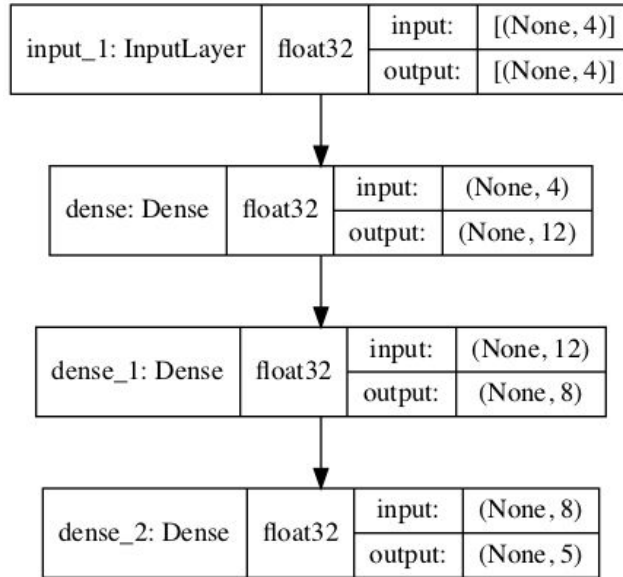
Credits: [What are Symbolic and Imperative APIs in TensorFlow 2.0?](#)

Two broad styles: Symbolic and Imperative

For a more rigorous treatment of these two styles, check out:

[What are Symbolic and Imperative APIs in TensorFlow 2.0?](#)

Models with `tf.keras.Sequential`



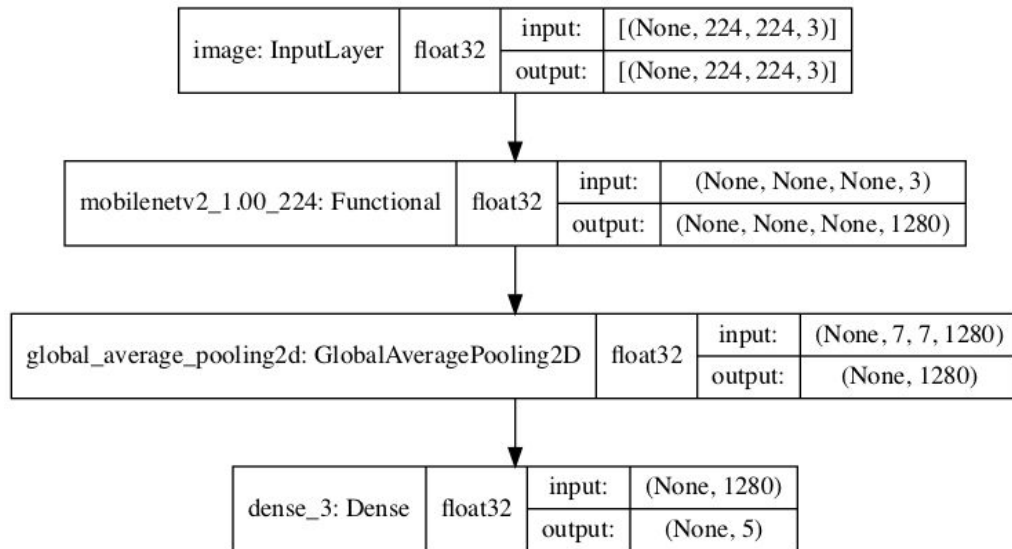
A sequential stack of layers

```
num_features = 4
```

```
num_classes = 5
```

```
sequential_model = keras.Sequential([  
    keras.layers.InputLayer(input_shape=(num_features,)),  
    keras.layers.Dense(12, activation="relu"),  
    keras.layers.Dense(8, activation="relu"),  
    keras.layers.Dense(num_classes, activation="softmax")  
], name="sequential_model")
```

Image classification model

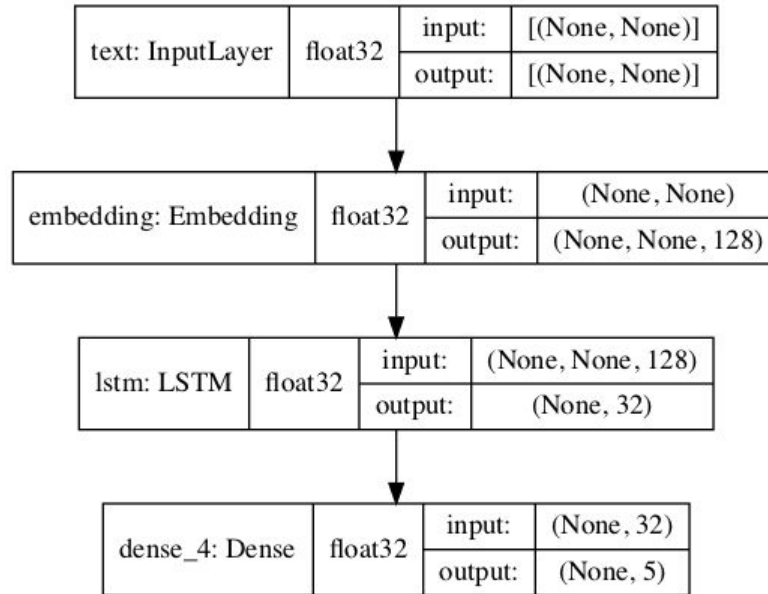


**Pretrained backbone +
Classification top**

```
mobilenet = keras.applications.MobileNetV2(weights="imagenet", include_top=False)
image_resolution = 224

custom_mobilenet_sequential = keras.Sequential(
    [
        keras.layers.InputLayer(
            input_shape=(image_resolution, image_resolution, 3), name="image"
        ),
        mobilenet,
        keras.layers.GlobalAveragePooling2D(),
        keras.layers.Dense(num_classes, activation="softmax"),
    ],
    name="custom_mobilenet_sequential",
)
```

Text classification model



A sequential stack of layers

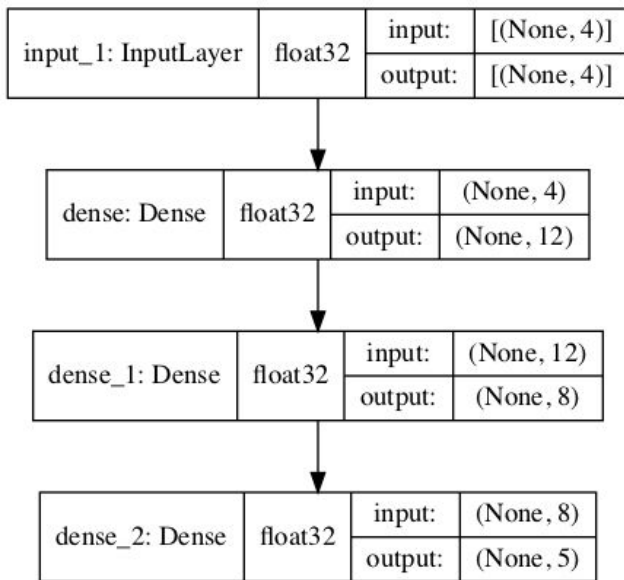
```
embedding_dim = 128
input_dim = 512
lstm_units = 32

text_model_sequential = keras.Sequential(
    [
        keras.layers.InputLayer(input_shape=(None,), name="text"),
        keras.layers.Embedding(input_dim=input_dim, output_dim=embedding_dim),
        keras.layers.LSTM(units=lstm_units),
        keras.layers.Dense(num_classes, activation="softmax"),
    ],
    name="text_model_sequential",
)
```

Models with the Functional API

- What if you need to return multiple outputs from a single model?
- What if your model needs to accept multiple inputs of multiple modalities (image, text, audio, etc.)?
- What if you need to add a skip connection in between the intermediate layer outputs?
- ...

Multi-output model with the Functional API

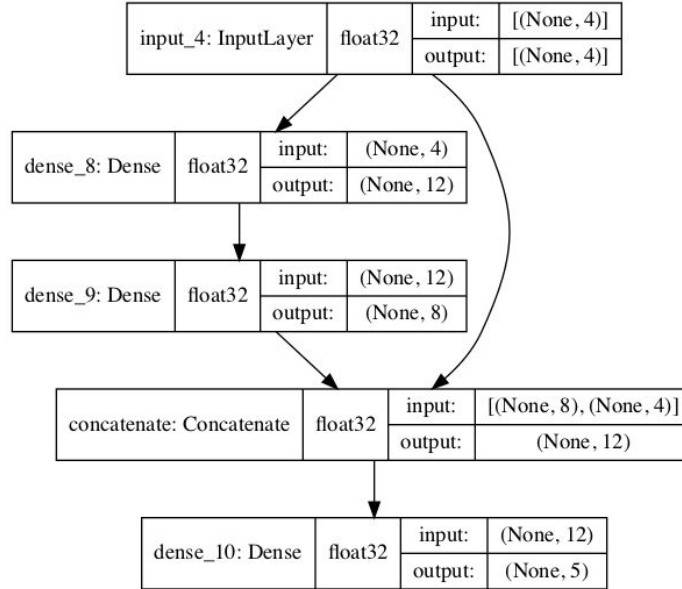


**Model computation graph
unchanged**

```
num_features = 4
num_classes = 5

inputs = keras.Input((num_features,))
x1 = keras.layers.Dense(12, activation="relu")(inputs)
x2 = keras.layers.Dense(8, activation="relu")(x1)
outputs = keras.layers.Dense(num_classes, activation="softmax")(x2)
functional_model_multi = keras.Model(
    inputs, outputs=[x1, x2, outputs], name="functional_model"
)
```

Model with combined intermediate outputs



**The computational flow is not
entirely sequential**

```
num_features = 4
```

```
num_classes = 5
```

```
inputs = keras.Input((num_features,))
```

```
x = keras.layers.Dense(12, activation="relu")(inputs)
```

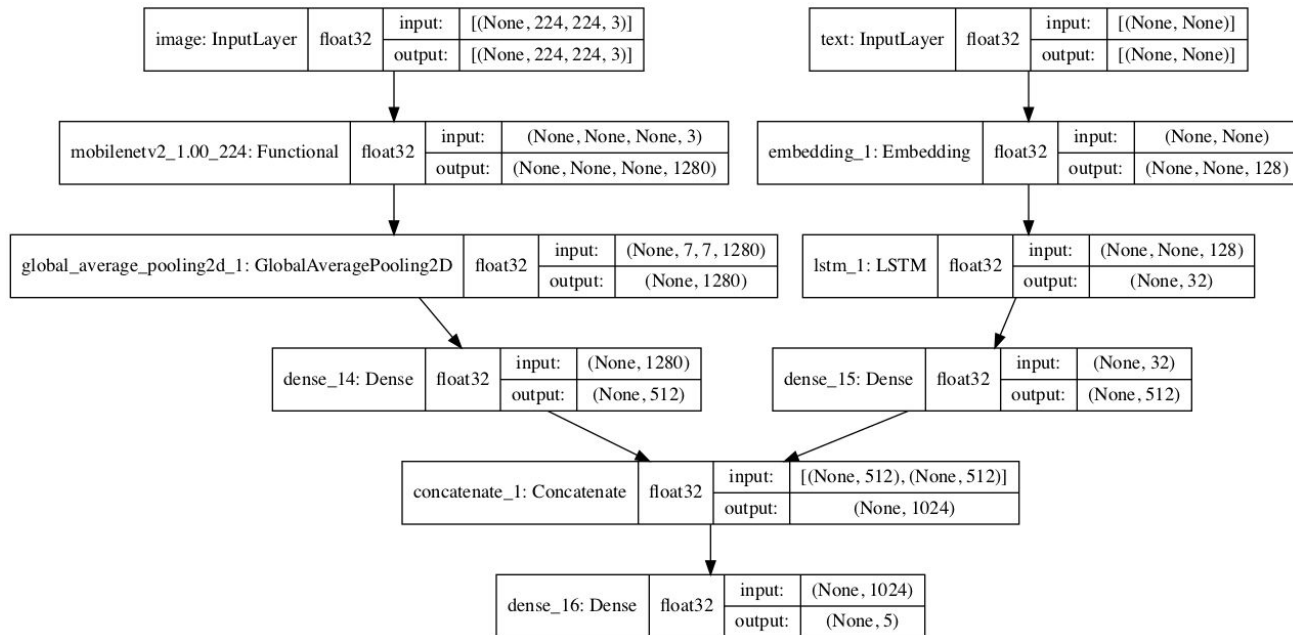
```
x = keras.layers.Dense(8, activation="relu")(x)
```

```
concatenated_features = keras.layers.Concatenate()([x, inputs]) ←
```

```
outputs = keras.layers.Dense(num_classes, activation="softmax")(concatenated_features)
```

```
functional_model = keras.Model(inputs, outputs, name="functional_model")
```

Model with multiple inputs



**A model accepting images and
texts as inputs**

```
# Inputs
```

```
image_inputs = keras.Input((image_resolution, image_resolution, 3), name="image")
```

```
text_inputs = keras.Input((None,), name="text")
```

```
# MobileNet model for images
```

```
image_model = sequential_image_model()
```

```
# Image representations
```

```
image_representations = image_model(image_inputs)
```

```
projected_image_representations = keras.layers.Dense(projection_dim, activation="relu")(
    image_representations
)
```

```
# Text representations
```

```
lstm_output = functional_text_model(text_inputs)
```

```
projected_text_representations = keras.layers.Dense(projection_dim, activation="relu")(
    lstm_output
)
```

```
...
```

...

```
# Concatenate image and text representations
```

```
concatenated_projections = keras.layers.Concatenate()(  
    [projected_image_representations, projected_text_representations]  
)
```

```
# Classification top
```

```
outputs = keras.layers.Dense(num_classes, activation="softmax")(concatenated_projections)  
multimodal_model = keras.Model([image_inputs, text_inputs], outputs, name="multimodal_model")
```

Models with subclassing from `keras.Model`

- Allows you to take more granular controls of what happens inside your model.
- Particularly useful if your model has a tree-like or recursive structure.


```
class ShallowMLP(keras.Model):

    def __init__(self, num_classes=num_classes, **kwargs):
        super().__init__(**kwargs)
        self.dense1 = keras.layers.Dense(12, activation="relu")
        self.dense2 = keras.layers.Dense(8, activation="relu")
        self.classification_layer = keras.layers.Dense(
            num_classes, activation="softmax"
        )

    def call(self, inputs):
        x = self.dense1(inputs)
        x = self.dense2(x)
        outputs = self.classification_layer(x)
        return outputs
```

Some gotchas

- No structured outputs when you call `summary()`.
- No concrete model computation graph with `plot_model()`.
- More during the code walkthrough and [here](#).

Customize the training behaviour

- Subclass from `keras.Model` overriding the `train_step()` method.
- Control the evaluation loop by overriding `test_step()`.

```
class Trainer(keras.Model):
    def train_step(self, data):
        x, y = data

        with tf.GradientTape() as tape:
            # Forward pass
            y_pred = self(x, training=True)
            # Compute loss
            loss = keras.losses.sparse_categorical_crossentropy(y, y_pred)

            # Compute gradients
            trainable_vars = self.trainable_variables
            gradients = tape.gradient(loss, trainable_vars)
            # Update weights
            self.optimizer.apply_gradients(zip(gradients, trainable_vars))
            # Compute gradient norm
            norm_tracker.update_state(tf.norm(gradients[0]))
            # Compute our own metrics
            loss_tracker.update_state(loss)
            accuracy_metric.update_state(y, y_pred)

        return {
            "loss": loss_tracker.result(),
            "accuracy": accuracy_metric.result(),
            "first_layer_weight_grad_norm": norm_tracker.result(),
        }
```

```
inputs = keras.Input((num_features,))
x = keras.layers.Dense(12, activation="relu")(inputs)
x = keras.layers.Dense(8, activation="relu")(x)
outputs = keras.layers.Dense(num_classes, activation="softmax")(x)

# Notice that we're using `Trainer` instead of `keras.Model`.
functional_model_custom = Trainer(inputs, outputs, name="functional_model_custom")
```

Now call `fit()`

In [37]:

```
random_labels = tf.experimental.numpy.random.randint(
    0, 5, size=(random_inputs.shape[0],)
)

functional_model_custom.compile(optimizer="adam")

functional_model_custom.fit(random_inputs, random_labels, epochs=3)
```

2022-09-12 10:20:50.069967: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes (registered 2)

Epoch 1/3

1/1 [=====] - 0s 222ms/step - loss: 1.5824 - accuracy: 0.2500 - first_layer_weight_grad_norm: 4.4891

Epoch 2/3


1/1 [=====] - 0s 3ms/step - loss: 1.5757 - accuracy: 0.2500 - first_layer_weight_grad_norm: 4.4810

Epoch 3/3


1/1 [=====] - 0s 2ms/step - loss: 1.5691 - accuracy: 0.2500 - first_layer_weight_grad_norm: 4.4733


Enough talk, show me the code!

bit.ly/models-colab



Questions?





Thank you!



Sayak Paul
ML Engineer at Carted
[@RisingSayak](#)

